



RTG π^3 Compact Course A3

Deep learning for Digital Pathology and Inverse Problems

Daniel Otero Baguer, Peter Maass

Center for Industrial Mathematics (ZeTeM)
University of Bremen

Bremen
22-24.04.2020

Outline

- 1 Introduction
- 2 Convolutional Neural Networks
- 3 Architectures
- 4 Data augmentation
- 5 Transfer learning



Section 1

Introduction



Sources

- *"Deep Learning for Vision Systems"*, Mohamed Elgendy
<https://livebook.manning.com/book/deep-learning-for-vision-systems/welcome/v-8>
- *"Convolutional Neural Networks for Visual Recognition"*
<http://cs231n.stanford.edu/>



Why Deep learning now?

- Five decades of research in machine learning
- CPUs/GPUs/storage developed for other purposes
- Lots of data from the internet
- Imagenet



Why Deep learning now?

- Five decades of research in machine learning
- CPUs/GPUs/storage developed for other purposes
- Lots of data from the internet
- Imagenet



Why Deep learning now?

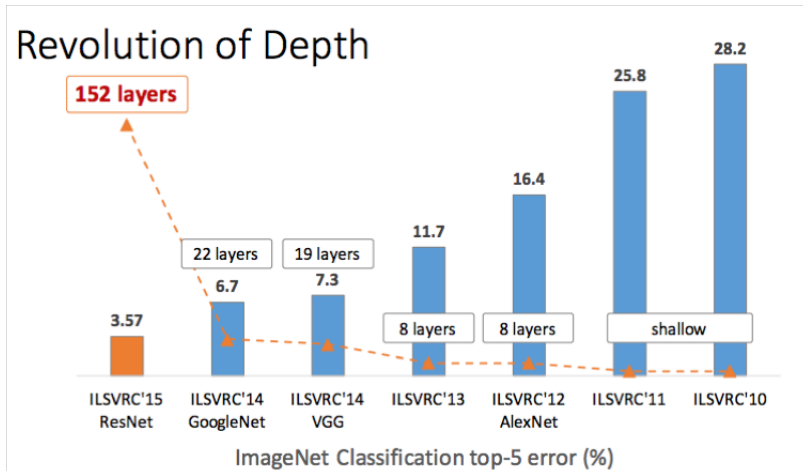
- Five decades of research in machine learning
- CPUs/GPUs/storage developed for other purposes
- Lots of data from the internet
- Imagenet

Why Deep learning now?

- Five decades of research in machine learning
- CPUs/GPUs/storage developed for other purposes
- Lots of data from the internet
- Imagenet



Why Deep Learning now?





Why Deep Learning now?

- Resources and efforts from large corporations make it **easier!**

Why Deep Learning now?

- Resources and efforts from large corporations make it **easier!**



theano



Microsoft
CNTK

PYTORCH



dmlc
mxnet



The learning task

Let's consider a simple regression problem where the observations are:

$$(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)}) \in X \times Y \quad (1)$$

Task: Infer a function f from the training data, which minimizes

$$R(f) = \int_{X \times Y} c(x, y, f(x)) dP(x, y) \quad (2)$$

where c is a loss function. A common choice is $c(x, y, f(x)) = \frac{1}{2} \|f(x) - y\|^2$.



The learning task

Empirical Risk

$$R(f) = \frac{1}{m} \sum_{i=1}^m c(x^{(i)}, y^{(i)}, f(x^{(i)})) \quad (3)$$

If we allow f to be any function that maps from X to Y , then we can minimize (3) but at the same time be very distant from the minimizer of (2).

$$f(x) = \begin{cases} y^{(i)} & x = x^{(i)} \\ 0 & \text{otherwise} \end{cases} \quad (4)$$

The learning task

No Free lunch theorem

If we do not make any assumptions on the class of functions where f belongs to, there is no chance to learn anything.

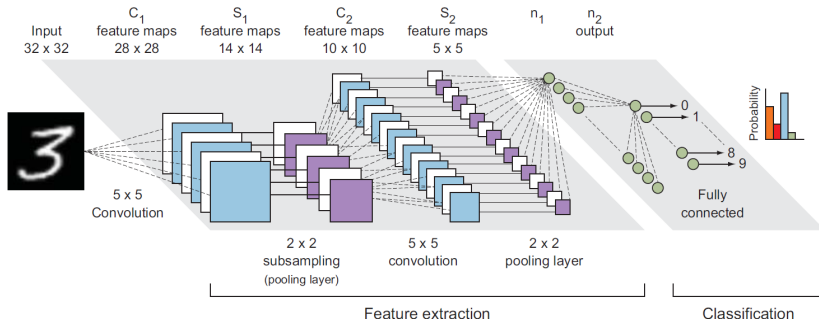




Section 2

Convolutional Neural Networks

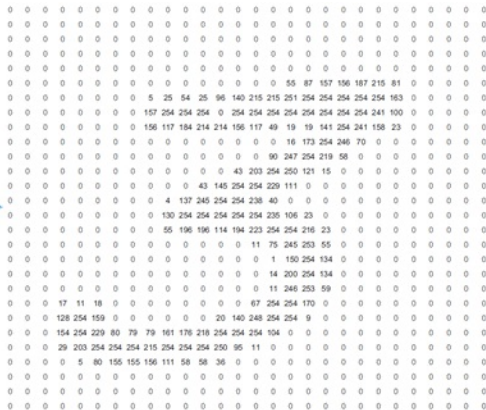
Convolutional Neural Network



- Input

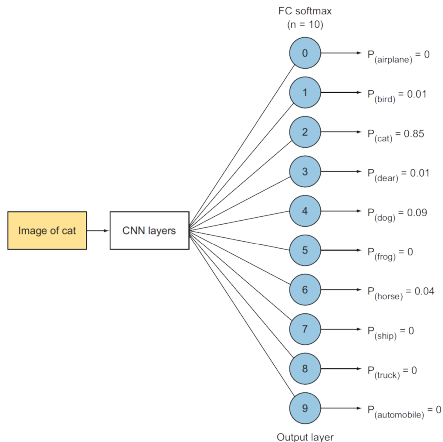


28 x 28
= 784 pixels




Layer types

■ Output



Layer types

■ Convolutions



Layer types

■ Convolutions

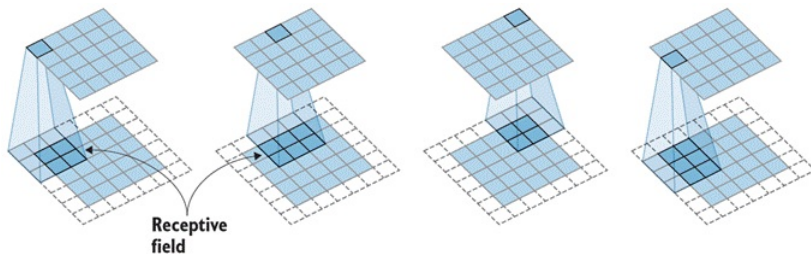


Figure: Receptive field

Layer types

■ Convolutions

Input image



Edge detection
kernel

*

0	-1	0
-1	4	-1
0	-1	0

=

Convolved image
(feature map)



Figure: Convolutional layers are inspired on standard image filtering

Layer types

■ Pooling

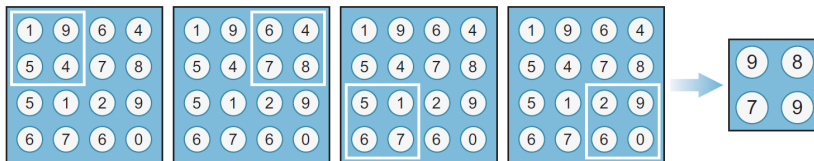
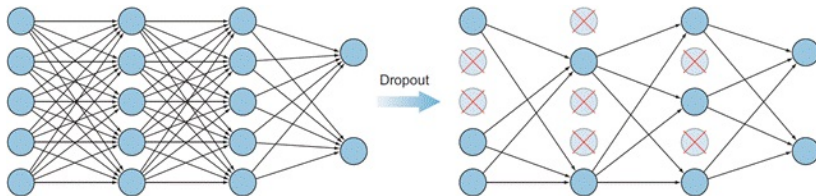


Figure: Max pooling example

Layer types

■ Dropout



Tensors




Figure: Tensors ¹

¹Image source: <https://towardsdatascience.com/understanding-pytorch-with-an-example-a-step-by-step-tutorial>




Section 3

Architectures

LeNet

- Introduced in 1998, by LeCun et al. in their paper “Gradient-Based Learning Applied to Document Recognition”
- 5 layers: 3 conv + 2 fully-connected
- 61706 parameters

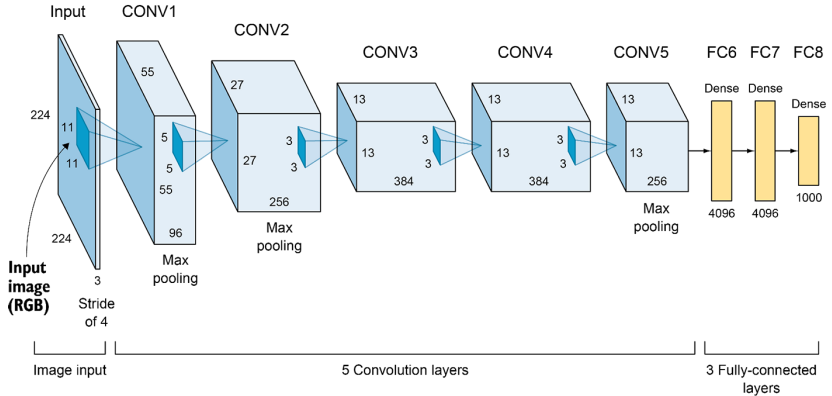




AlexNet

- Winner of the ILSVRC image classification competition in 2012.
- Introduced by Alex Krizhevsky, Geoffrey Hinton and Ilya Sutskever in their paper “ImageNet Classification with Deep Convolutional Neural Networks”
- 8 layers: 5 conv + 3 fully-connected
- 60 million parameters

AlexNet





How deep can we go?

- Very deep networks are able to represent very complex functions
- The network can learn features at many different levels of abstraction

Vanishing gradients

- By the chain rule, the derivatives of each layer are multiplied down the network
- Gradient decreases exponentially as we propagate down to the initial layers
- First hidden layers are learning much slower than later hidden layers

How deep can we go?




Figure: Sigmoid activation

How deep can we go?




Figure: Vanishing gradient effect. First layers train much slower.



How deep can we go?

Solutions:

- Use ReLu activations
- Normalization layers
- Residual Networks

How deep can we go?

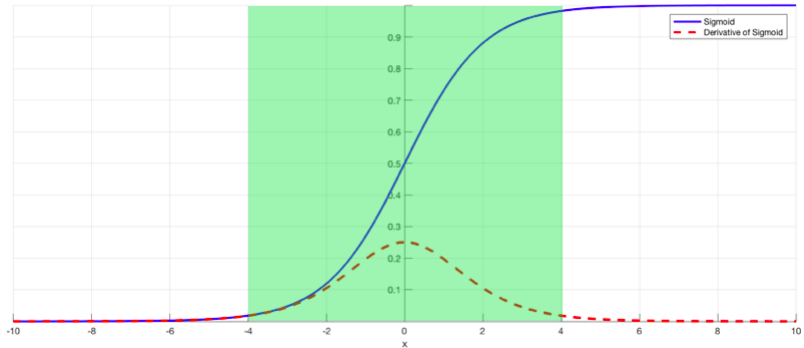


Figure: Sigmoid with restricted inputs

ResNet

- Introduces shortcut (skip-connection) that allows the gradient to directly back-propagate to earlier layers
- Allows the layer to learn an identity function (will perform at least as good as the previous layer)




Figure: ResNet's skip connection

ResNet

- From 18 to 152 layers
- With 152 layers: ≈ 60 million parameters




Figure: Residual block

ResNet




Figure: ResNet



Section 4

Data augmentation

Data augmentation

Main idea: Generate additional images based on existing ones

- Random flip
- Random erasing
- Random sized crop
- Random perspective
- Color jitter

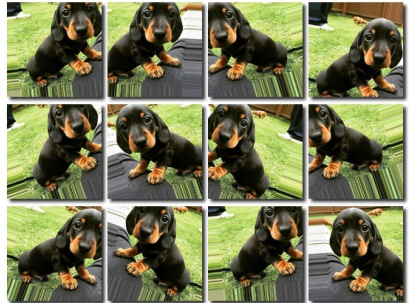


Figure: Data augmentation



Data augmentation

Insights

- The epoch size does not change
- We get different randomly transformed samples every epoch (dynamic dataset)
- Helps to avoid overfitting

Advise

- Applying heavy augmentations unnecessarily can result in poor accuracy
- Do not use data augmentation in validation or test set

Data augmentation

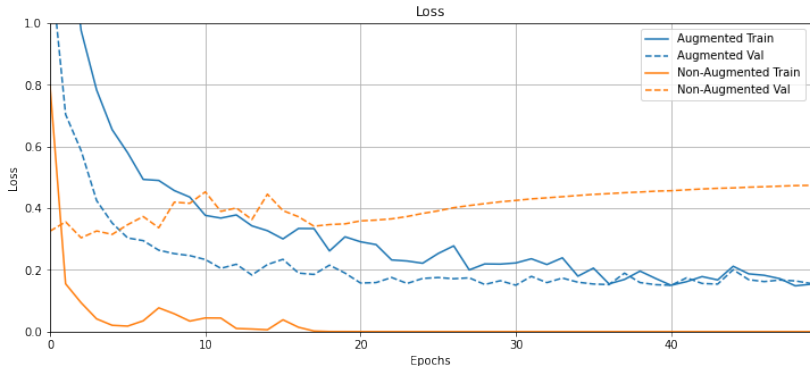


Figure: Non-augmented model vs augmented model ².

²https://www.tensorflow.org/tutorials/images/data_augmentation



Section 5

Transfer learning

Transfer learning

Common phenomena on CNNs trained on natural images:



Figure: First layer learn features similar to Gabor filters and color blobs. Each of the 96 filters shown here is of size $[11 \times 11 \times 3]$,

Claim: Such a layer is not specific to a particular dataset or task



Transfer learning

Main idea: Use a pre-trained network

Use cases

- Fixed feature extractor: Take a pre-trained network and train only the last fully-connected layer
- Fine-tuning: Fine-tune the weights of the pre-trained network by continuing the back-propagation
 - Use small learning rate
 - Keep some of the earlier layers fixed (due to overfitting concerns)



When and how?

	Small dataset	Large dataset
Similar	Fixed feature extractor + Linear Classifier	Fine-tuning
Very different	Keep early layers fixed + Linear classifier	Pre-trained as initialization



Practical advice

- Constraints from pre-trained models
- Small learning rate for fine-tuning (avoid distorting pre-trained parameters too quickly or too much)